

Dependency-order and preference-forest DP differential compression of the gapped suffix array

1 Setup and preliminaries

1.1 Gapped k -mers and DisLex

Fix a binary shape S of *span* s and *weight* w (number of care positions). On a DNA text $T[0..n)$, every start position p yields a gapped k -mer name $c = \text{name}_S(T, p)$. For the running example we take

$$T = (\text{GCCTTTAAAG})^3, \quad |T| = n = 30, \quad S = \# . \#, \quad s = 3, \quad w = 2.$$

The **DisLex** transform (Horton) lays these names out grouped by residue class $p \bmod s$, producing an integer string (the *lertext*) of length $m = n + 1 = 31$. Consecutive symbols inside a residue class are exactly the successive windows $p, p + s, p + 2s, \dots$ — i.e. the gapped suffix starting at p . The suffix array of the lextext is the **gapped suffix array** (GSA); an LCP array is computed in symbol space (Kasai). Writing $\text{sa}[r]$ for the lextext rank- r suffix and $\text{orig}(r)$ for its original text start, the first symbol of that suffix is the gapped k -mer at $\text{orig}(r)$.

1.2 LCP intervals I_c

Because the GSA is ordered by gapped suffixes, all ranks whose first symbol equals a fixed name c form a contiguous interval

$$I_c = [\ell_c, r_c) \subseteq \{0, \dots, m - 1\}.$$

The multiset $\{\text{orig}(r) : r \in I_c\}$ is precisely the sorted list of text positions of c . On the example, the intervals with $|I_c| > 1$ are:

c	SA range I_c	$ I_c $	text positions $\text{orig}(\cdot)$ (rank order)
AA	[2, 5)	3	26, 16, 6
AG	[5, 10)	5	27, 18, 8, 17, 7
CT	[10, 16)	6	22, 21, 12, 2, 11, 1
GC	[17, 22)	5	19, 9, 20, 10, 0
TA	[22, 28)	6	25, 24, 15, 5, 14, 4
TT	[28, 31)	3	23, 13, 3

(Singletons $\$ \$$, $\text{A\$}$, $\text{G\$}$ are never differentially compressed.)

1.3 Differential links

Let $x = \text{sa}[r]$ be a lextext position. Shifting left by a symbols in the lextext corresponds to walking a windows backward along the same residue class. Whenever the alignment condition

$$\text{lex2orig}[x - a] = \text{orig}(r) - a \cdot s$$

holds, the suffix at $x - a$ is a valid **left context** of depth a . Prepending a symbols preserves relative order among suffixes and shifts original positions by $+a \cdot s$.

Consequently, if a contiguous SA run $J = [\sigma, \tau)$ has LCP depth at least $a + 1$ and the symbol at depth a equals c , then mapping each kept source rank $u \in J$ by

$$\text{orig}(u) \longmapsto \text{orig}(u) + a \cdot s$$

recovers, bijectively and in order, a subset of I_c . We call (J, a) a **differential source** of coverage $\text{cov} = |J'|$ for those covered ranks of I_c .

1.4 What is stored: C , H_{offset} , H_{rest}

Compression drops covered ranks from the stored position array. The kept original positions, still in GSA rank order, form the compressed array C . For each name c the hash table stores at most two entries:

- $H_{\text{offset}} = (\text{pos}, a, \text{num})$: recover num positions as $C[\text{pos} + t] + a \cdot s$ for $t = 0, \dots, \text{num} - 1$;
- $H_{\text{rest}} = (\text{pos}, \text{num})$: num positions stored directly in C .

Ranks used as sources of an accepted link are **pinned** (their pin-count increases) so that later decisions cannot remove them while dependents still need them. The objective of source selection is to minimise $|C|$.

2 The preference relation

For each interval I_c with $|I_c| > 1$, the implementation enumerates all viable sources for $a = 1, \dots, a_{\text{max}}$ (default $a_{\text{max}} = 8$):

1. For every rank $r \in I_c$, compute the predecessor lertext coordinate $y = x - a$ if the alignment condition above holds; collect pairs $(\text{rank}(y), r)$.
2. Sort by source rank and group into maximal contiguous runs where consecutive source ranks differ by 1 and $\text{LCP} \geq a + 1$.
3. Retain a run only if it is **disjoint** from I_c and has coverage at least 2.

Among all such candidates, the **preferred** candidate for I_c maximises coverage, breaking ties by preferring larger a (deeper LCP). Availability (pins / already-removed ranks) is *ignored* when building preferences: preference is a global, static notion of “who would like whom as differential source.”

Directionality is one-sided: sources always lie to the left in the lertext (earlier windows on the same residue class). There is no “forward” differential link.

3 The DepOrder algorithm

DepOrder has three conceptual stages.

3.1 Preference DAG

For every interval with a preferred candidate of coverage ≥ 2 , draw a directed edge

$$\text{source name} \longrightarrow \text{target name}$$

meaning “the target prefers a differential source inside the source’s interval.” Self-loops are discarded. The resulting digraph may contain cycles.

On the example the preferred candidates are:

target	prefers	a	cov	source SA range	covered ranks in target
AA	GC	2	3	[19, 22)	{2, 3, 4}
AG	CT	2	5	[11, 16)	{5, 6, 7, 8, 9}
CT	GC	4	2	[20, 22)	{10, 12}
GC	CT	3	2	[14, 16)	{19, 20}
TA	CT	1	6	[10, 16)	{22, ..., 27}
TT	GC	1	3	[19, 22)	{28, 29, 30}




Figure 1: Preference DAG for $T = (\text{GCCTTTAAAG})^3$, $S = \#.\#$. An edge $w \rightarrow c$ means that c 's preferred differential source lies inside I_w . Hubs CT and GC form a mutual-preference cycle; leaves AA, TT, AG, TA are high-coverage dependents.

Hubs. Both CT and GC are preferred sources for three names each. There is a mutual preference cycle $\text{CT} \leftrightarrow \text{GC}$. The high-value dependents are the leaves AA, TT (via GC) and AG, TA (via CT): together they offer coverage $3 + 3 + 5 + 6 = 17$, while the hub–hub edges themselves offer only coverage 2 each.

3.2 Phase I — reverse topological acceptance (sinks first)

A Kahn topological order is computed on the preference DAG; remaining nodes in any strongly connected component (here the CT–GC cycle and its waiting dependents) are appended by decreasing interval size. **Acceptance proceeds in reverse topological order:** dependents / sinks before sources.

For each name in that order, the algorithm recomputes the best candidate *under current availability* (source run must be fully kept; covered ranks must be free of pins and not already removed) and, if coverage ≥ 2 , accepts it: pin the source ranks, mark covered ranks removed.

Why sinks first helps hubs. Accepting a leaf that points into a hub pins the hub's ranks before the hub itself is considered. A hub that would like to compress *itself* away (using the other hub) then finds its own ranks already pinned and is forced to **skip**, remaining available as a source for all dependents. Size-first greedy does the opposite: it lets a large hub claim a weak slice of another hub early, spoiling high-coverage leaf links.

3.3 Phase II — unpin / retarget with transitive add

After Phase I, a fixed-point pass repeatedly tries, for intervals in decreasing size order, candidates of *strictly larger* coverage than the current acceptance — temporarily ignoring pin constraints on covered ranks. Accepting such a candidate may require **retargeting** dependents that currently pin those ranks: if a dependent D 's source run is a contiguous sub-run of the new candidate's covered ranks, D is rewritten to point at the corresponding sub-run of the new source with

$$a'_D = a_D + a_{\text{new}}.$$

The attempt is transactional (snapshot / rollback) and kept only if the global number of kept ranks decreases, or coverage for that name improves. Phase II is where former sources can themselves be compressed once their dependents have been lifted onto a deeper source.

4 Worked example: GCCTTTAAAG $\times 3$ with #.#

Parameters: $n = 30$, $s = 3$, $m = 31$, $a_{\max} = 8$. Distinct names: 9.

4.1 Phase I on the example

Because of the CT–GC cycle, Kahn’s order first drains the indegree-0 singletons, then appends the cyclic component by size. Reverse-topo acceptance among compressible names is

TT, AA, GC, AG, TA, CT.

step	interval	decision	detail
I.1	TT	ACCEPT	$\leftarrow$ GC[19, 22), $a = 1$, cov=3. Pins GC ranks 19–21.
I.2	AA	ACCEPT	$\leftarrow$ GC[19, 22), $a = 2$, cov=3. Same hub pins reinforced.
I.3	GC	SKIP	Preferred source CT[14, 16) would cover GC ranks {19, 20}, but those ranks are already pinned by TT/AA. Availability blocks self-compression of the hub.
I.4	AG	ACCEPT	$\leftarrow$ CT[11, 16), $a = 2$, cov=5. Pins CT ranks 11–15.
I.5	TA	ACCEPT	$\leftarrow$ CT[10, 16), $a = 1$, cov=6. Pins full CT hub.
I.6	CT	SKIP	Preferred GC[20, 22) is pinned; covered CT ranks are themselves pinned as sources. Hub preserved.

After Phase I: **kept** =14, four accepted offsets (AA, AG, TA, TT). Phase II finds no improving retarget on this instance, so the final size remains $|C| = 14$.

Explicitly, C stores the three singletons plus the entire hubs CT and GC:

$$C = [30, 28, 22, 21, 12, 2, 11, 1, 29, 19, 9, 20, 10, 0].$$

Hash table (offset / rest):

c	H_{offset}	H_{rest}
\$\$	—	(0, 1)
A\$	—	(1, 1)
AA	(11, 2, 3)	—
AG	(3, 2, 5)	—
CT	—	(2, 6)
G\$	—	(8, 1)
GC	—	(9, 5)
TA	(2, 1, 6)	—
TT	(11, 1, 3)	—

All 17 non-hub, non-singleton positions are recovered from the two hubs: $3 + 5 + 6 + 3 = 17$, and $31 - 17 = 14$.

4.2 Contrast: size-order greedy ($|C| = 22$)

Greedy processes intervals by decreasing $|I_c|$, with availability, and pins sources permanently:

1. **CT** (size 6) accepts GC[20, 22), $a = 4$, cov only **2** — a weak use of two GC hub ranks to drop two of six CT positions.
2. **TA** (size 6) is then limited to the same spoiled GC slice: cov only **2**.
3. **AG** (size 5) **skips** — its preferred CT source is no longer fully available.
4. **GC** falls back to TT[29, 31), $a = 2$, cov **2**.
5. **AA** still accepts GC[19, 22), cov 3.
6. **TT skips**.

Net removal: $2 + 2 + 2 + 3 = 9$ ranks, hence $|C| = 22$. The structural mistake is clear: by letting the hub CT claim a low-coverage fragment of GC first, greedy destroys the high-coverage leaf edges (AG cov 5, TA cov 6, TT cov 3) that DepOrder’s sinks-first order realises in full.

algorithm	$ C $	fraction of $m = 31$
Greedy	22	71.0%
DepOrder	14	45.2%

(On this instance Greedy-DFS reaches 15 — better than greedy, still one above DepOrder.)

5 Complexity and when DepOrder helps

Phase I is a single pass over a topological (or size-refined) order, each step re-enumerating candidates for one interval — inexpensive relative to GSA construction. Phase II may repeat enumeration, snapshotting, and retarget attempts over $O(N)$ rounds ($N = \#$ intervals) and is the asymptotic bottleneck on large instances; on the present toy example it is a no-op.

DepOrder helps precisely when the preference graph has **high fan-out hubs** whose own preferred self-compression edges are low-coverage, while many dependents offer large coverage into those hubs. Sinks-first pinning protects the hubs; Phase II’s transitive add is reserved for the harder case where a hub should later be compressed *through* a deeper source after dependents have been retargeted. The GCCTTTAAAG $\times 3 / \# . \#$ instance is the minimal didactic witness of the hub-protection phenomenon: CT and GC must stay in C , and every other multi-occurrence name collapses onto them, yielding $|C| = 14$ against greedy’s 22.

6 The Tree-DP algorithm (preference-forest DP)

Tree-DP attacks the same objective — minimise $|C|$ — by an exact dynamic program on a *preference forest* extracted from the same static preferences, followed by the shared Phase II pass.

6.1 From preference DAG to preference forest

Recall that each interval I_c with $|I_c| > 1$ has a preferred differential candidate (maximising coverage, then a). Tree-DP retains a candidate for name c only if

$$|I_c| > 2$$

(singletons and size-2 intervals are left to a later leftover pass). A directed preference $w \rightarrow c$ (“ c prefers a source inside I_w ”) becomes a **forest edge** only when its coverage satisfies

$$\text{cov} \geq 3.$$

Self-loops are dropped, and an edge is skipped if it would close a cycle in the forest under construction. The resulting graph is therefore a forest of rooted trees: each non-root c has a unique parent $\text{parent}(c) = w$, the preferred source of c .

Weak mutual preferences of coverage 2 (such as CT $\leftrightarrow$ GC on the running example) are filtered out by the $\text{cov} \geq 3$ threshold, so cycles in the full preference DAG typically dissolve into independent high-coverage stars.

6.2 KEEP versus COMPRESS

Write $\text{isize}(v) = |I_v|$. For a node v whose preferred source is currently available ($\text{src_ok} = \text{true}$), Tree-DP chooses between two local decisions:

- **KEEP** v : store all $\text{isize}(v)$ ranks of I_v . Every child w may then treat v as an available source:

$$\text{cost}_{\text{keep}}(v) = \text{isize}(v) + \sum_{w \text{ child of } v} \text{DP}(w, \text{src_ok} = \text{true}).$$

- **COMPRESS** v : accept the preferred candidate of coverage $\text{cov}(v)$, paying only the uncovered ranks $\text{isize}(v) - \text{cov}(v)$. Children lose v as a source:

$$\text{cost}_{\text{comp}}(v) = \text{isize}(v) - \text{cov}(v) + \sum_{w \text{ child of } v} \text{DP}(w, \text{src_ok} = \text{false}).$$

The recurrence returns the cheaper option (KEEP on a tie). If $\text{src_ok} = \text{false}$, COMPRESS is forbidden and the node is forced to KEEP. Roots are handled identically, except that COMPRESS is allowed only when the preferred source lies *outside* the root’s preference subtree (so the source is not itself a forest descendant of the root).

The DP is exact for the $|C|$ contribution of all forest nodes under this KEEP / COMPRESS model. Intervals that never entered the forest (or remain uncompressed after the DP choice) are finished by a size-ordered availability greedy; then the shared Phase II unpin / retarget fixed point of DepOrder is run unchanged.

6.3 Worked example on the preference forest

On $T = (\text{GCCTTTAAAG})^3$, $S = \#.\#$, the $\text{cov} \geq 3$ filter retains exactly four edges and yields two stars:

parent (source)	child (target)	a	cov
GC	AA	2	3
GC	TT	1	3
CT	AG	2	5
CT	TA	1	6

Leaves under an available hub. With $\text{src_ok} = \text{true}$, each leaf prefers COMPRESS: e.g. $\text{cost}_{\text{comp}}(\text{AA}) = 3 - 3 = 0$ (versus KEEP 3), and likewise TT, AG, TA all compress to cost 0.




Figure 2: Preference forest for the running example: two stars $GC \rightarrow \{AA, TT\}$ and $CT \rightarrow \{AG, TA\}$. The weak hub–hub edges of coverage 2 do not enter the forest.

Root GC.

$$\begin{aligned} \text{cost}_{\text{keep}}(GC) &= 5 + 0 + 0 = 5, \\ \text{cost}_{\text{comp}}(GC) &= (5 - 2) + \underbrace{3}_{\text{AA KEEP}} + \underbrace{3}_{\text{TT KEEP}} = 9 \end{aligned}$$

(compressing GC via CT would deny the children their hub, forcing them to KEEP). Since $5 < 9$, the DP **KEEPS** GC.

Root CT.

$$\begin{aligned} \text{cost}_{\text{keep}}(CT) &= 6 + 0 + 0 = 6, \\ \text{cost}_{\text{comp}}(CT) &= (6 - 2) + 5 + 6 = 15. \end{aligned}$$

Again KEEP wins. Materialisation accepts the four leaf compressions (deepest first), leftover greedy finds nothing further, and Phase II is a no-op — exactly as under DepOrder — yielding

$$|C| = 14.$$

6.4 Contrast with DepOrder

Both algorithms protect the same hubs on this instance and reach the same $|C| = 14$. They differ in mechanism:

	DepOrder	Tree-DP
Graph model	preference DAG (cycles allowed)	preference forest: nodes with $ I_c > 2$; edges only if $\text{cov} \geq 3$ (cycle-breaking)
Phase I	reverse-topo acceptance with live availability / pins (sinks first)	exact KEEP / COMPRESS DP per tree; materialise deepest-first
Hub–hub cov 2	present in the DAG; sinks-first pinning forces hubs to SKIP	omitted from the forest; DP compares $\text{cost}_{\text{keep}}$ (5 / 6) to the costly COMPRESS alternative
Phase II	unpin / retarget / transitive add	same shared Phase II

DepOrder reasons operationally (who pins whom, in which order). Tree-DP reasons globally on a simplified forest: the numeric comparison $\text{cost}_{\text{keep}} < \text{cost}_{\text{comp}}$ is precisely the statement that keeping a hub is cheaper than compressing it at the expense of its high-coverage dependents — the same structural insight DepOrder realises by sinks-first pinning.

7 Summary

On the didactic instance $\text{GCCTTTAAAG} \times 3 / \# . \#$, both DepOrder and Tree-DP store the hubs CT and GC in full, differentially recover AA, TT, AG, and TA from them, and finish at $|C| = 14$ (against size-order greedy’s 22). Phase II is common to both and idle here; it matters when a kept hub should later be lifted onto a deeper source after dependents have been retargeted. Tree-DP’s filters $|I_c| > 2$ and $\text{cov} \geq 3$ turn the cyclic preference DAG into a clean two-star forest on which KEEP versus COMPRESS is decided by dynamic programming.